

OCAD · ONE COMPANY A DAY
SEMICONDUCTOR


NVIDIA

NVIDIA Corporation

NASDAQ: NVDA · United States · Semiconductors

USD 225.07

Closing price, 25 September 2026 (last completed trading session)



Executive Summary

Business

NVIDIA designs graphics and accelerated-computing processors and sells them, together with a proprietary software and networking stack, primarily to hyperscale cloud operators, enterprises and AI infrastructure builders. Data Center compute and networking is now the company's overwhelming revenue engine, with Gaming, Professional Visualization and Automotive as smaller, historically core franchises. Revenue is concentrated among a small number of hyperscale and AI-infrastructure customers and is increasingly exposed to United States export policy toward China.

Financial Condition

Fiscal 2026 (ended 25 January 2026) revenue was USD 215.9 billion, up 65 percent year over year, with GAAP net income of USD 120.1 billion (55.6 percent net margin) and free cash flow of USD 96.7 billion. For the trailing twelve months ended 26 July 2026, revenue reached USD 303.0 billion and net income USD 192.9 billion. The balance sheet held USD 99.4 billion of cash and marketable securities against USD 33.4 billion of debt as of 26 July 2026, a net cash position of approximately USD 66.0 billion, after the company issued USD 25.0 billion of senior unsecured notes in the June 2026 quarter.

Market Position

NVIDIA's principal competitors in AI accelerators and graphics are Advanced Micro Devices and, indirectly, the custom-silicon programs of its own largest customers; Broadcom and Marvell compete for custom AI accelerator design; Intel and Qualcomm compete in adjacent compute and edge silicon. Its principal advantage is the CUDA software ecosystem and a full-stack hardware-networking-software platform that raises switching costs for developers and infrastructure operators. Its principal weakness is dependence on a single external foundry, Taiwan Semiconductor Manufacturing Company, for advanced-node wafer capacity, and a revenue base concentrated among a handful of large buyers.

Valuation

At the 25 September 2026 close of USD 225.07 per share and 24.147 billion shares outstanding, market capitalization was approximately USD 5.43 trillion and enterprise value approximately USD 5.37 trillion. The stock traded at 28.5 times trailing GAAP diluted earnings per share, 24.9 times forward consensus earnings, and approximately 17.7 times trailing twelve-month revenue on an enterprise-value basis.

Investment Framework

ATTRACTIVE

Data Center revenue has compounded at triple-digit rates for three consecutive fiscal years, gross margin sits above 70 percent, and the balance sheet carries a net cash position despite an aggressive capital-return programme.

CONCERNING

Four customers accounted for 61 percent of quarterly revenue as of the most recently disclosed period, versus 36 percent a year earlier, and reported quarterly revenue growth has decelerated from triple digits toward the 30 to 40 percent range implied by the company's own guidance.

MUST GO RIGHT	Hyperscale and sovereign AI infrastructure capital spending must continue to grow at a pace that absorbs NVIDIA's expanding accelerator, networking and system output without a material pricing or share give-back to custom silicon.
COULD BREAK THE THESIS	A slowdown in AI infrastructure capital expenditure by the small number of customers that dominate Data Center revenue, a faster-than-expected shift toward in-house accelerators at those same customers, or renewed export restrictions that foreclose additional addressable markets.
PRICED IN	The analysis in Section B3 indicates that, on the base-case discounted cash flow assumptions set out in Section C5, the current share price requires revenue to compound at roughly 60 percent a year for five years at a sustained free-cash-flow margin near 42 percent, a materially more optimistic path than this report's base case.

Investment Profile

BUSINESS QUALITY	High. Dominant share of AI accelerator compute, high recurring engagement through CUDA and its software stack, gross margin above 70 percent.
FINANCIAL QUALITY	High. Net cash balance sheet, free cash flow conversion above 40 percent of revenue, cash generation materially exceeds reported capital expenditure.
COMPETITIVE POSITION	Strong but narrowing. Leading merchant-silicon position in AI accelerators; increasing competitive pressure from hyperscaler in-house silicon and from Broadcom-designed custom accelerators.
GROWTH PROFILE	Decelerating from an exceptional base. Revenue growth has fallen from above 100 percent to a still-high double-digit to low-triple-digit range as the comparison base scales.
BALANCE-SHEET RISK	Low. Net cash position of approximately USD 66.0 billion even after a USD 25.0 billion bond issuance and continued buybacks and dividends.
VALUATION	Demanding. Trailing and forward multiples sit well above the semiconductor peer group and imply growth and margin assumptions above this report's base case, as quantified in Section B3.
MARKET AND TECHNICAL PROFILE	High realized volatility (40 to 50 percent annualized) and a beta near 2.2 versus the S&P 500; the stock trades above both its 50-day and 200-day moving averages with a neutral 14-day RSI near 44.

A. About the Business

A1. Business Environment

NVIDIA operates in the semiconductor and accelerated-computing industry, specifically the design of graphics processing units (GPUs), data-center accelerators, networking silicon and the software layers that run on them. The company is fabless: it designs chips and subcontracts manufacturing to third-party foundries, principally Taiwan Semiconductor Manufacturing Company (TSMC), and outsources packaging, testing and much of final system assembly to contract manufacturers. The industry value chain therefore runs from electronic design automation tools and intellectual-property licensing, through wafer fabrication and advanced packaging, to system integration by original equipment and design manufacturers, and finally to end deployment inside hyperscale cloud data centers, enterprise data centers, and consumer devices.

Revenue pools in this industry are concentrated at two points: leading-edge logic fabrication capacity (captured largely by TSMC) and the highest-performance accelerator and networking silicon designs (captured by NVIDIA and a small number of competitors). Cost pools are concentrated in research and development, foundry wafer and packaging costs, and, for buyers, in the capital expenditure required to build and power data centers. Capital intensity for the fabless design model itself is comparatively low relative to integrated device manufacturers, but the industry NVIDIA sells into, hyperscale AI infrastructure, is extremely capital intensive, with the largest cloud operators collectively guiding to several hundred billion dollars of annual capital expenditure. The industry is cyclical at the semiconductor-equipment and memory level and has historically exhibited multi-year up and down cycles in capacity and pricing; the current AI-accelerator sub-segment has grown so quickly from a small base that a full multi-year cycle has not yet been observed.

The industry is technologically dependent on process-node advancement (TSMC's leading-edge nodes), on advanced packaging (CoWoS and successor technologies), and on high-bandwidth memory supplied by a small number of memory manufacturers, creating supply-chain concentration risk shared across the industry rather than unique to NVIDIA. Geographically, design and headquarters functions are concentrated in the United States, manufacturing is concentrated in Taiwan and, to a lesser extent, South Korea, and the largest pools of demand are concentrated among United States hyperscale cloud operators alongside a growing set of sovereign and enterprise AI infrastructure buyers.

INDUSTRY STAGE CLASSIFICATION

The AI-accelerator segment of the semiconductor industry is best classified as **structurally transforming and still growing at a triple-digit rate off an enlarging base**, nested inside a broader semiconductor industry that is mature and cyclical. The growth segment is structurally transforming because computing demand is shifting from general-purpose central processing units toward purpose-built accelerators for training and inference of large-scale AI models, a shift that is reshaping capital allocation across the entire technology sector. It is not yet appropriately described as mature, because reported quarter-over-quarter growth for the industry leader remained in the high-teens percent range through mid-2026 and forward guidance continued to call for further sequential growth. It carries speculative characteristics at the level of public-market valuation, discussed in Section B3, even though the underlying revenue and cash flow growth are real and sourced.

A2. What Does the Company Do?

NVIDIA earns money by designing semiconductor chips, and the systems and software built around them, and selling or licensing that output to companies that need to run large-scale computing workloads. The clearest way to understand the business is by end market:

Data Center. The dominant and fastest-growing part of the business. NVIDIA sells GPU accelerators (the Blackwell architecture in volume during fiscal 2026 and 2027, with the successor Rubin architecture disclosed for future generations), central processing units designed in-house (the Grace and successor Vera CPUs), networking silicon and switches (the


Spectrum-X Ethernet and NVLink/InfiniBand product lines from the Mellanox acquisition), and the BlueField data-processing unit for storage, networking and security offload. These are sold both as individual components to system builders and increasingly as fully integrated rack-scale systems, to cloud service providers, large consumer internet companies, enterprises and government or sovereign AI programmes.

Gaming. GeForce GPUs for personal computers, sold through add-in-board partners and system builders to retail and gaming consumers; the company's original business and still a multi-billion dollar, profitable franchise, though now a small fraction of total revenue.

Professional Visualization. RTX-branded workstation graphics and the Omniverse simulation and digital-twin software platform, sold to design, media, manufacturing and engineering customers.

Automotive and Robotics. The DRIVE platform for autonomous and assisted driving compute, and the Jetson and Isaac platforms for robotics, sold to automakers, their tier-one suppliers and robotics developers, both as current-generation hardware revenue and as a call option on longer-term autonomy adoption.

Revenue mix by end market, fiscal 2026 (full year)




Units: percent of fiscal 2026 total revenue of USD 215.9 billion. Source: NVIDIA fourth-quarter and fiscal 2026 results release.

Beginning with the quarter ended 26 July 2026, NVIDIA's public revenue disclosure was simplified to two market platforms, Data Center and Edge Computing, with the latter aggregating Gaming, Professional Visualization and Automotive and Robotics; the two internal reportable accounting segments remain Compute and Networking and Graphics. This report uses the four-category breakdown where the fiscal 2026 annual disclosure provides it, and the two-category breakdown for the most recent quarter, and states in each case which convention is used.

Unlike a pure hardware vendor, a material share of NVIDIA's durable advantage sits in software: the CUDA parallel-computing platform, associated libraries (cuDNN, TensorRT, NIM microservices) and the Omniverse and AI Enterprise

software suites. These lower the cost of building on NVIDIA hardware and raise the cost of moving workloads to a competitor's silicon once a developer's code, tooling and internal expertise are built around CUDA.

OPERATING MODEL: HOW A DOLLAR ENTERS THE COMPANY AND WHERE IT GOES



Sourced: NVIDIA fourth-quarter and fiscal 2026 results release; NVIDIA second-quarter fiscal 2027 results release and CFO commentary (see Section G).

The most economically important links in this chain are the Data Center revenue box, which now determines the trajectory of the entire company, and the foundry relationship embedded in the cost box, since NVIDIA's own operating margin and delivery schedule are directly gated by wafer and advanced-packaging capacity it does not own or control.

A3. Business and Relationship Network

NVIDIA sits at the center of a concentrated ecosystem of suppliers, customers and partners that is unusually consequential for a company of its size, because both its manufacturing input and a large share of its revenue are concentrated among a small number of counterparties.



The TSMC relationship matters because it is a single point of failure for leading-edge output: any disruption to Taiwanese fabrication or advanced packaging capacity would directly constrain NVIDIA's ability to ship its most profitable products, and no comparably capable second-source foundry currently exists at the same process node. The customer-concentration relationships matter because a small number of buyers now determine the trajectory of company-wide revenue; a capital-expenditure pause at any one of the top four accounts would have a disproportionate effect on reported growth. The Intel

equity relationship is notable because it links NVIDIA, as an investor and technology partner, to a competitor that is simultaneously a potential alternative source of advanced x86 CPU and foundry capacity for the United States semiconductor ecosystem; its economic significance to NVIDIA's own income statement is presently secondary to the operating relationships above.

A4. Competition

NVIDIA's most economically relevant competitors are selected on the basis of who competes for the same AI-infrastructure capital budgets, not merely companies that share a semiconductor industry classification. TSMC is treated as an ecosystem dependency (a critical supplier) rather than a competitor, since it does not compete with NVIDIA for end-customer accelerator spending.

COMPETITIVE COMPARISON

Company	What it sells	Target customer	Scale (TTM revenue)	Relative position vs NVIDIA
Advanced Micro Devices (AMD)	Instinct AI accelerators, EPYC server CPUs, gaming and embedded GPUs	Hyperscale and enterprise data centers, PC OEMs	USD 41.3 billion	Direct merchant-silicon competitor in AI accelerators and CPUs; smaller scale, materially lower operating margin, thinner software ecosystem than CUDA
Broadcom Inc.	Custom AI accelerator (ASIC) design, networking silicon, enterprise software	Hyperscale cloud operators designing in-house chips	USD 89.1 billion	Indirect but strategically important competitor: enables the largest cloud customers to reduce reliance on merchant GPUs through custom silicon programmes
Intel Corporation	x86 CPUs, Gaudi AI accelerators, foundry services	PC OEMs, enterprise and cloud data centers	USD 57.0 billion	Weaker current competitor in AI accelerators; potential long-term alternative in CPU and (via its foundry business) manufacturing capacity; NVIDIA is also a minority investor
Qualcomm Incorporated	Mobile and edge system-on-chip processors, emerging data-center inference accelerators	Mobile device makers; emerging edge and inference AI infrastructure buyers	Not retrieved	Indirect competitor concentrated in edge and mobile inference rather than large-scale training, a smaller overlap with NVIDIA's core Data Center franchise
Hyperscaler in-house silicon (Google TPU, Amazon Trainium/Inferentia, Microsoft Maia)	Custom AI accelerators developed for internal cloud use	The hyperscaler's own internal workloads and, increasingly, external cloud customers	Not separately disclosed	Emerging competitor and the single largest structural risk to NVIDIA's Data Center franchise, since these accelerators substitute directly for merchant GPU purchases inside the customers that matter most to NVIDIA's revenue

TTM revenue for AMD, Broadcom and Intel from Alpha Vantage company data as of late September 2026 (see Section G). Qualcomm and hyperscaler in-house programmes are not separately broken out by their parent companies in a directly comparable revenue line and are marked Not Retrieved rather than estimated.

Each competitor sells what they sell to whom they sell it for a distinct reason. AMD competes head-on for the same merchant-silicon AI accelerator budget with a similar go-to-market model, making it the cleanest like-for-like comparison.

Broadcom does not sell branded accelerators to end customers; it sells design services and silicon intellectual property that allow NVIDIA's own largest customers to build competing chips, which makes it a competitor for capital rather than for a specific product sale. Intel competes across CPUs and, more speculatively, accelerators, and simultaneously participates in NVIDIA's investor base, an unusual dual relationship. Hyperscaler in-house silicon is the least visible but most consequential competitive force, because it is deployed by the same companies that generate the bulk of NVIDIA's revenue.

A5. Competitive Differences

Advantages. NVIDIA's most durable advantage is the CUDA software ecosystem: more than a decade of accumulated libraries, developer tooling, trained personnel and optimized AI framework support that raises the switching cost of moving a production AI workload to a competing architecture, independent of any single generation of hardware. This is reinforced by scale advantages in R&D spending, which fund a roughly annual architecture cadence (Hopper, Blackwell, and the disclosed successor Rubin) that competitors have struggled to match, and by a full-stack systems and networking offering (GPUs, CPUs, NVLink and Spectrum-X networking, and system software) that lets NVIDIA sell an integrated rack rather than a component, capturing more of the total system value and increasing switching costs at the infrastructure level. Brand and ecosystem trust with both developers and capital markets further support pricing power, evidenced by gross margins sustained above 70 percent even as volumes have grown more than tenfold in three years.

Disadvantages. The company is fully dependent on a single external foundry for leading-edge manufacturing capacity, giving it no direct control over the most capital-intensive and geopolitically sensitive part of its supply chain. Revenue is concentrated among a small number of customers, several of whom are simultaneously NVIDIA's largest addressable competitive threat through their own in-house silicon programmes, an unusual customer-competitor overlap. The business is also directly exposed to trade policy: the fiscal 2026 10-K disclosed a USD 4.5 billion inventory charge tied to the H20 product and stated that, as of fiscal year end, NVIDIA was effectively foreclosed from China's data-center computing market, illustrating that a large addressable market can be removed by policy decisions outside the company's control.

Distinguishing temporary from structural: the current scale of Data Center revenue growth (triple-digit annual rates) is not structural and should be expected to moderate as the comparison base enlarges, which the company's own guidance already reflects. The CUDA ecosystem advantage, by contrast, is structural, because it does not depend on any single product cycle and compounds with every additional developer and enterprise deployment built on it.

MOAT ASSESSMENT Strong. The moat is anchored in software switching costs (CUDA) and a full-stack systems advantage, partially offset by single-source foundry dependence and customer concentration that is itself a source of competitive threat. The moat is not rated Very Strong because the largest customers have both the capital and the technical capability to internalize an increasing share of accelerator demand over a multi-year horizon.

A6. Industry Trend

Demand for AI-accelerated computing has grown from a specialized workload into the primary driver of data-center capital spending among the largest cloud operators, evidenced by NVIDIA's own Data Center revenue compounding from roughly USD 15 billion in fiscal 2023 to USD 193.7 billion in fiscal 2026. The near-term driver is continued buildout of large language model training and inference capacity, agentic AI workloads and enterprise AI adoption; NVIDIA's own commentary at the Q2 fiscal 2027 results release characterized demand as being constrained by supply rather than by orders, with the company describing "compute is revenue" and citing capacity commitments extending toward calendar 2032.

Supply-side constraints, principally advanced-node wafer and packaging capacity at TSMC and high-bandwidth memory supply, have been the binding limit on how quickly NVIDIA can convert demand into recognized revenue, rather than a shortage of customer orders. Pricing trends have favored NVIDIA: gross margin has remained above 70 percent through a

period of more than tenfold volume growth, indicating pricing power rather than the margin erosion that typically accompanies rapid volume scaling in semiconductors. Regulatory and technological change is significant and two-sided: export restrictions on advanced accelerators to China have removed a portion of the addressable market (Sourced: fiscal 2026 10-K), while government and sovereign AI infrastructure programmes in other jurisdictions have partially offset this by becoming new, discrete demand pools.

The industry is best classified as **growing and structurally transforming**, with cyclical risk building beneath the surface: the customer base funding this growth is narrow, capital-intensive, and ultimately dependent on the funding, revenue and profitability of AI applications built on top of this infrastructure. A reversal of the current trend would most plausibly originate from a slowdown in hyperscale capital expenditure guidance, evidence of AI-infrastructure overbuild relative to realized application revenue, or a faster-than-anticipated transition of large customers toward in-house silicon.

A7. Management

NAMED EXECUTIVE OFFICERS

Name	Position	Background
Jen-Hsun (Jensen) Huang	Co-Founder, President and Chief Executive Officer	Co-founded NVIDIA in 1993 and has served as CEO throughout the company's history; previously held engineering and management roles at LSI Logic and Advanced Micro Devices
Colette M. Kress	Executive Vice President and Chief Financial Officer	Joined NVIDIA in 2013; previously chief financial officer of the business technology and operations group at Cisco Systems and held finance roles at Microsoft and Texas Instruments
Ajay K. Puri	Executive Vice President, Worldwide Field Operations	Joined NVIDIA in 2005; previously held senior sales and marketing roles at Sun Microsystems
Debora Shoquist	Executive Vice President, Operations	Joined NVIDIA in 2007; prior operations leadership roles at technology and optical-components companies
Timothy S. Teter	Executive Vice President, General Counsel and Secretary	Joined NVIDIA in 2017; previously a patent litigation partner at Cooley LLP

Sourced: NVIDIA fiscal 2026 Form 10-K, executive officer disclosures (see Section G).

Jensen Huang's continuous tenure as founder-CEO across more than three decades is unusual in large-capitalization technology and is a relevant factor in evaluating strategic continuity: the pivot from graphics into general-purpose accelerated computing and, subsequently, into AI infrastructure was a multi-year strategic bet initiated well before the current AI infrastructure buildout was visible to the broader market, and it has been financially vindicated by the revenue and margin trajectory documented in Section C1. Capital allocation under the current management team has combined heavy R&D and supply-commitment reinvestment with substantial, growing shareholder returns: USD 40.1 billion of buybacks in fiscal 2026, a materially enlarged quarterly dividend introduced in calendar 2026 (from a nominal rate to USD 0.25 per share, payable 1 October 2026), and the company's first large-scale bond issuance (USD 25.0 billion, June 2026 quarter), which funds continued infrastructure investment while preserving net cash. This combination is consistent with capital allocation aligned with shareholder interests rather than empire-building; the H20 charge disclosed in the fiscal 2026 10-K is a reminder that management's China assumptions can be overtaken by policy events outside its control.

B. Evaluation


B1. Financial and Operating Metrics

PROFITABILITY AND RETURNS

Metric	FY2022	FY2023	FY2024	FY2025	FY2026	TTM Jul-26
Gross margin	64.9%	56.9%	72.7%	75.0%	71.1%	74.7%
Operating margin	37.3%	20.7%	54.1%	62.4%	60.4%	65.2%
Net margin	36.2%	16.2%	48.8%	55.8%	55.6%	63.7%
Return on equity (calculated)	36.6%	19.8%	69.2%	91.9%	76.3%	84.2%
Return on assets (calculated)	n.a.	n.a.	n.a.	n.a.	58.1%	60.2%


Reasoned/Calculated: margins are gross profit, operating income and net income each divided by revenue for the corresponding period; return on equity is net income divided by period-end shareholders' equity; return on assets is net income divided by period-end total assets. Underlying revenue, income and balance-sheet figures are Sourced from NVIDIA's fiscal 2026 results release and its Q2 fiscal 2027 Form 10-Q (see Section G).

Revenue and year-over-year growth, FY2022-FY2026 and TTM




Units: US\$ billions (bars, left axis); percent (line, right axis). Fiscal years ended late January. Source: NVIDIA results releases, FY2022-FY2026 and Q2 FY2027.

Gross, operating and net margin, FY2022-FY2026 and TTM



Units: percent of revenue. Reasoned/Calculated from Sourced income-statement data. Source: NVIDIA results releases.

Return on equity, FY2022-FY2026



Units: percent. Reasoned/Calculated as net income divided by period-end shareholders' equity. Source: NVIDIA results releases.

Return on invested capital for fiscal 2026, calculated as after-tax operating income (GAAP operating income of USD 130.4 billion, taxed at the 15.1 percent effective rate disclosed for fiscal 2026) divided by invested capital (total debt plus shareholders' equity less cash and equivalents, USD 155.2 billion at fiscal year end), was approximately **71.3 percent**; on the same methodology using trailing-twelve-month operating income and the 26 July 2026 balance sheet, ROIC was approximately **69.9 percent**. Both figures are Reasoned/Calculated and sit far above any plausible estimate of NVIDIA's cost of capital (Section C5), consistent with a business earning substantial economic rents on invested capital, a hallmark of a durable competitive advantage rather than a commodity manufacturing return profile.

GROWTH (CALCULATED COMPOUND ANNUAL GROWTH RATES)

Metric	3-year CAGR (FY2023-FY2026)	4-year CAGR (FY2022-FY2026)
Revenue	100.2%	68.3%

Metric	3-year CAGR (FY2023-FY2026)	4-year CAGR (FY2022-FY2026)
Diluted EPS	n.a.	88.3%
Free cash flow	n.a.	85.7%

Reasoned/Calculated from Sourced annual figures in Section C1 and C3. Growth of this magnitude reflects the company's transition from a graphics-centered business to the dominant supplier of AI data-center infrastructure and should not be extrapolated mechanically; see Section C4 for forward scenarios.

EFFICIENCY AND CASH CONVERSION

Metric	FY2026	Latest available
Asset turnover (revenue / total assets)	1.04x	0.95x (TTM)
Operating cash flow / net income	85.6%	n.a.
Free cash flow margin	44.8%	39.3% (H1 FY2027)
Capital expenditure / revenue	2.8%	n.a.
Free cash flow yield (on current market capitalization)	1.8% (FY2026 FCF basis)	2.6% (H1 FY2027 FCF, annualized)

Reasoned/Calculated. Operating cash flow of USD 102.7 billion against net income of USD 120.1 billion in fiscal 2026 means cash generation ran below reported earnings for the year, a normal pattern when receivables and inventory expand alongside triple-digit revenue growth, and worth monitoring rather than dismissing; it is not, by itself, evidence of low earnings quality. Reported capital expenditure of USD 6.0 billion in fiscal 2026 understates NVIDIA's true exposure to the AI infrastructure buildout, because the great majority of data-center capital spending is undertaken by NVIDIA's customers, not by NVIDIA itself; NVIDIA's own capital intensity remains low because it is a fabless designer that outsources manufacturing.

Leverage metrics are of limited relevance because NVIDIA carries a net cash position: total debt of USD 33.4 billion against cash and marketable securities of USD 99.4 billion as of 26 July 2026, for net cash of approximately USD 66.0 billion, equivalent to roughly negative 0.3 times trailing EBITDA. Gross debt to trailing EBITDA was approximately 0.17 times. Interest coverage is Not Retrieved as a discrete disclosed line item, but is not a source of near-term concern given the net cash position; the more relevant leverage question, addressed in Section B2, is whether the newly issued USD 25.0 billion of senior notes marks the start of a more debt-funded capital programme as NVIDIA's own infrastructure ambitions (including data-center and energy investments disclosed alongside the bond issuance) expand.

B2. Risks

STRUCTURED RISK REGISTER, RANKED BY PROBABILITY × IMPACT

Risk	Category	Probability	Impact
Customer concentration: a small number of hyperscale buyers set the pace of revenue growth	Business	High	High
Accelerated shift by top customers toward in-house AI accelerators	Competitive / Technological	Medium	High
Further export-control tightening removing additional addressable markets	Regulatory	Medium	High
Single-foundry dependence on TSMC for advanced-node and packaging capacity	Business / Supply Chain	Low to Medium	High
Valuation multiple compression if growth decelerates faster than the market expects	Valuation	Medium	High

Risk	Category	Probability	Impact
AI infrastructure capital expenditure overbuild relative to realized application revenue at customers	Macro / Business	Low to Medium	High
Product-cycle execution risk on annual architecture cadence (Blackwell to Rubin transition)	Technological	Low to Medium	Medium
Rising debt issuance changing the balance-sheet risk profile if sustained	Financial	Low	Medium
Cybersecurity and intellectual-property protection across a complex global supply chain	Technological	Low to Medium	Medium
Key-person dependence on a single, long-tenured founder-CEO	Business / Governance	Low	Medium

THREE RISKS MOST CAPABLE OF PERMANENTLY IMPAIRING SHAREHOLDER VALUE

1. A durable, structural shift by the largest cloud customers toward in-house AI accelerators, which would simultaneously shrink NVIDIA's addressable market and strengthen the negotiating position of its remaining hyperscale buyers.
2. A sudden or sustained disruption to TSMC's advanced-node or packaging capacity, whether from geopolitical, natural, or operational causes, for which no comparable alternative supplier currently exists.
3. A sharp deceleration in AI-infrastructure capital spending driven by disappointing returns on AI applications at NVIDIA's largest customers, which would simultaneously compress reported growth and the growth expectations embedded in the current valuation.

Customer concentration is quantifiable and current: the most recently disclosed data (quarter ended October 2025) showed four direct customers accounting for 61 percent of quarterly revenue, each individually above 10 percent, versus three customers together accounting for 36 percent a year earlier. This is both a business risk (revenue is dependent on continued spending decisions by very few counterparties) and, notably, a competitive risk, because those same counterparties are the entities best resourced to develop competing in-house silicon.

B3. Valuation


CURRENT VALUATION MULTIPLES

Multiple	NVDA	Basis
Price / trailing diluted EPS	28.5x	Price USD 225.07 / TTM diluted EPS USD 7.90
Price / forward consensus EPS	24.9x	Data-provider consensus forward estimate
EV / trailing revenue	17.7x	EV USD 5,368.8bn / TTM revenue USD 303.0bn
EV / trailing EBITDA	23x to 27x	Cross-checked range: 23.2x per third-party data provider, 26.7x on this report's own EV calculation
Price / free cash flow (FY2026 basis)	56.2x	Market cap USD 5,434.8bn / FY2026 FCF USD 96.7bn
Price / free cash flow (annualized H1 FY2027 run-rate)	38.9x	Market cap / (H1 FY2027 FCF of USD 69.9bn annualized)
Dividend yield	approx. 0.1%	Trailing dividend per share blending the pre- and post-increase quarterly rate

Reasoned/Calculated using Sourced price, share count and financial-statement inputs (Section G). EV/EBITDA is shown as a range because this report's own enterprise-value calculation and a third-party data provider's EBITDA estimate produce different results; both are disclosed for transparency rather than presenting a single false-precision figure. Market data as of 25-26 September 2026.

HISTORICAL VALUATION

Price / trailing diluted EPS at each fiscal year end, FY2022-FY2026, and current




Reasoned/Calculated: fiscal year-end closing share price divided by that fiscal year's trailing GAAP diluted EPS. Source: monthly closing prices and NVIDIA results releases.

The trailing price-to-earnings multiple has compressed from roughly 39 times at the fiscal 2026 year end to 28.5 times currently, even as the share price rose, because trailing earnings grew faster than the share price over the same period. This is the opposite of multiple expansion and indicates that, on a pure trailing-earnings basis, the market has not become more optimistic about NVIDIA in valuation terms over the past eight months; it has simply been correct that earnings would catch up.


PEER VALUATION

Market capitalization, peer comparison



Units: US\$ trillions. Source: Alpha Vantage company data and stockanalysis.com, late September 2026.

Trailing price / earnings, peer comparison



Intel shown as not meaningful (n.m.) due to negative trailing GAAP earnings. Source: as above.

PEER VALUATION AND PROFITABILITY SNAPSHOT

Company	TTM revenue (USD bn)	Net margin	Trailing P/E	EV/EBITDA
NVIDIA	303.0	63.7%	28.5x	23.2x
Broadcom	89.1	42.9%	44.9x	33.9x
Taiwan Semiconductor Mfg.	139.6	49.9%	28.9x	19.5x
Advanced Micro Devices	41.3	15.6%	161.7x	92.8x
Intel	57.0	(19.8%)	n.m.	182.0x

Sourced: Alpha Vantage company overview data (NVDA, AMD, AVGO, INTC) and stockanalysis.com (TSM), as of the week of 21 September 2026 (see Section G).

NVIDIA trades at a materially lower trailing multiple than Broadcom and at a level close to TSMC despite substantially higher profitability and growth than either, and at a small fraction of AMD's multiple, which is distorted by AMD's much lower net margin. On profitability-adjusted terms, NVIDIA does not appear to be the most expensive name in this peer set; the more relevant question, addressed below, is what growth the absolute multiple embeds.

WHAT THE CURRENT PRICE APPEARS TO EMBED

Using the discounted cash flow framework in Section C5 (weighted average cost of capital 15.7 percent, terminal growth 3.5 percent), matching the current enterprise value of approximately USD 5.37 trillion requires either (i) revenue compounding at approximately 60 percent a year for five straight years from the fiscal 2026 base, reaching roughly USD 2.3 trillion of revenue by the fifth forecast year, while sustaining a free cash flow margin near 42 percent, or (ii) a free cash flow margin near 90 percent sustained across this report's base-case revenue path, a level with no precedent in large-scale semiconductor history. Both are Reasoned, mechanical outputs of the DCF framework, not predictions; they are presented so the reader can judge for themselves whether the growth path required is Conservative, Reasonable, Optimistic or Extremely Optimistic. This report characterizes the growth path required to justify the current price as **Optimistic to Extremely Optimistic**: it requires the current exceptional growth phase to persist for materially longer, and at a materially higher sustained margin, than this report's base case assumes, though it is not without precedent given the company's actual trailing three-year growth rate has already exceeded 100 percent per year.

B4. Capital Market and Stock Analysis


PRICE HISTORY AND RETURNS (TO 25 SEPTEMBER 2026 CLOSE OF USD 225.07)

Period	Level / return
1-day change	+0.22%
1-month (25-Aug-26 to 25-Sep-26)	+5.6%
3-month (25-Jun-26 to 25-Sep-26)	+15.0%
6-month (approx., from 31-Mar-26)	+29.4%
1-year	+20.9%
3-year	+419.0%
5-year	+991.2%

Period	Level / return
52-week range	USD 163.90 - USD 236.00


Sourced/Reasoned: computed from monthly and daily adjusted closing prices (see Section G); 52-week range from Alpha Vantage company data. All returns are split-adjusted for the June 2024 ten-for-one stock split.

NVDA share price, five-year (split-adjusted monthly closes)



Units: US\$ per share, month-end close, split-adjusted. Source: compiled monthly closing-price data, September 2021-September 2026.

NVDA versus S&P 500, indexed total return



Both series rebased to 100 at 30 September 2021. S&P 500 proxied by the SPY exchange-traded fund, including dividends. Source: compiled monthly adjusted closing-price data.

VOLATILITY AND DRAWDOWN

Metric	Value
3-month daily volatility, annualized	40.5%

Metric	Value
5-year monthly volatility, annualized	49.7%
Beta versus S&P 500 (5-year monthly, calculated)	2.22
Beta versus S&P 500 (third-party data provider)	2.22
Trailing 3-month maximum drawdown (daily closes)	(19.4%)
Calendar 2026 year-to-date maximum drawdown (month-end closes)	(8.8%)
5-year maximum drawdown (month-end closes, Nov-2021 to Sep-2022)	(62.8%)

Reasoned/Calculated from Sourced daily and monthly closing-price series (see Section G). The independently calculated beta of 2.22 matches the third-party data provider figure exactly, corroborating both.

NVIDIA's realized volatility, at roughly 40 to 50 percent annualized, and its beta near 2.2, place it well above the broad market and above most large-capitalization semiconductor peers, consistent with a stock whose valuation is sensitive to shifts in growth expectations rather than merely to broad market moves. The 5-year maximum drawdown of 62.8 percent, concentrated in the November 2021 to September 2022 broad technology and rate-driven de-rating, is a reminder that a stock trading on optimistic growth assumptions can re-rate sharply lower even when the underlying business continues to grow, as NVIDIA's own revenue did throughout that period.

TECHNICAL STRUCTURE

TECHNICAL INDICATORS (AS OF 25 SEPTEMBER 2026)

Indicator	Value	Reading
50-day moving average	USD 215.62	Price trading above the 50-day average
200-day moving average	USD 199.28	Price trading above the 200-day average; 50-day above 200-day (a "golden cross" structure)
14-day Relative Strength Index (RSI)	43.7	Neutral; neither overbought (above 70) nor oversold (below 30)

Moving averages Sourced from third-party data provider; RSI Reasoned/Calculated by the analyst from the daily closing-price series in Section G.

The technical structure, price above both moving averages with the shorter average above the longer one, is consistent with an intact medium-term uptrend, while a neutral RSI indicates the stock is neither exhausted to the upside nor oversold following the correction earlier in 2026. This technical reading describes present market positioning only; it says nothing about whether the fundamental valuation discussed in Section B3 is or is not justified, and the two should not be conflated.

C. Financial Statements

All figures in this section are Reported (as filed) unless labeled Derived, Adjusted or Estimated. NVIDIA reports under United States GAAP with a fiscal year ending on the last Sunday of January; fiscal 2026 accordingly covers the period from late January 2025 to 25 January 2026, and is referenced throughout as FY2026. All dollar figures are in United States dollars.


C1. Income Statement

CONSOLIDATED INCOME STATEMENT SUMMARY, FY2022-FY2026 (US\$ MILLIONS EXCEPT PER-SHARE DATA)

Line item	FY2022	FY2023	FY2024	FY2025	FY2026
Revenue	26,914	26,974	60,922	130,497	215,938
Revenue growth (YoY)	n.a.	0.2%	125.9%	114.2%	65.5%
Gross profit	17,475	15,356	44,301	97,858	153,463
Gross margin	64.9%	56.9%	72.7%	75.0%	71.1%
Operating income	10,041	5,577	32,972	81,453	130,387
Operating margin	37.3%	20.7%	54.1%	62.4%	60.4%
Net income	9,752	4,368	29,760	72,880	120,067
Net margin	36.2%	16.2%	48.8%	55.8%	55.6%
Diluted EPS (US\$)	0.39	0.17	1.19	2.94	4.90

Sourced: NVIDIA fourth-quarter and fiscal-year results releases, FY2022-FY2026 (SEC-filed exhibits), cross-checked against stockanalysis.com. All EPS figures restated on a post-split basis following the June 2024 ten-for-one stock split. Revenue growth and margin rows are Reasoned/Calculated from the reported lines above them.

Revenue and year-over-year growth (repeated from Section B1 for reference)



Units: US\$ billions and percent. Source: NVIDIA results releases.

Revenue growth in FY2024 through FY2026 was overwhelmingly a Data Center and volume story rather than a pricing or acquisition story: Data Center revenue rose from a small base to USD 193.7 billion in FY2026 on unit and system volume driven by successive GPU architecture launches (Hopper, then Blackwell), not from price increases on existing products or from acquired revenue. Operating leverage was strongly positive through FY2025, as operating margin expanded from 37.3 percent (FY2022) to a peak of 62.4 percent (FY2025), before contracting to 60.4 percent in FY2026; the FY2026 contraction was driven primarily by the USD 4.5 billion H20 inventory and purchase-obligation charge tied to Chinese export restrictions (Sourced: FY2026 10-K), a one-time, policy-driven item rather than a structural cost or pricing deterioration. Excluding that charge, underlying operating margin would have remained closer to the FY2025 level, consistent with continued, if moderating, operating leverage.


Gross margin followed a similar pattern: it fell in FY2023 (56.9 percent) during an inventory correction unrelated to the current AI cycle, recovered sharply through FY2025 (75.0 percent) as the higher-margin Data Center mix scaled, and gave back roughly four points in FY2026 (71.1 percent) principally on the same China-related charge; the trailing-twelve-month gross margin of 74.7 percent through July 2026 indicates the underlying margin has substantially recovered. Margins are best described as **structurally expanding with the Data Center mix shift, subject to episodic compression from trade-policy items that are outside management's control.**

C2. Balance Sheet

BALANCE SHEET SUMMARY, FY2024-FY2026 AND LATEST QUARTER (US\$ MILLIONS)

Line item	FY2024 (Jan-24)	FY2025 (Jan-25)	FY2026 (Jan-26)	Q2 FY2027 (Jul-26)
Cash and cash equivalents	7,280	8,589	10,605	22,443
Marketable securities	n.a.	n.a.	51,951	76,926
Goodwill (goodwill and intangibles combined for FY2024-FY2025)	5,500	6,000	20,832	21,125
Total assets	65,728	111,601	206,803	320,272
Total debt	10,828	9,982	8,468	33,366
Total liabilities	22,750	32,274	49,510	91,288
Shareholders' equity	42,978	79,327	157,293	228,984
Net cash / (net debt)	n.a.	n.a.	54,088	66,003

Sourced: NVIDIA fiscal 2026 Form 10-K and Q2 fiscal 2027 Form 10-Q balance sheets, and the FY2026 results release (see Section G). FY2024 and FY2025 cash figures and total debt from stockanalysis.com (cross-checked against reported equity and total assets); marketable securities not separately available for those years in the summarized source and shown as not available (n.a.) rather than estimated. Net cash / (net debt) is Reasoned/Calculated as cash and marketable securities less total debt.

Total debt and cash & investments

Units: US\$ billions. Debt series FY2022-FY2026 from stockanalysis.com (includes finance-lease liabilities); Q2 FY2027 debt and all cash & investments figures from NVIDIA's own reported balance sheet. FY2022-FY2025 cash & investments not shown as a consistent series was not available; see table above.


The balance sheet strengthened dramatically over the period shown, with shareholders' equity growing more than fivefold from FY2024 to FY2026 as retained earnings accumulated from record profitability. Goodwill and intangible assets rose sharply in FY2026, to USD 20.8 billion from USD 6.0 billion in FY2025, reflecting acquisition activity during the year; this remains a modest 10 percent of total assets and does not, on its own, constitute a red flag. The most significant balance-sheet development of the period is the USD 25.0 billion senior unsecured note issuance completed in the June 2026 quarter, which lifted long-term debt from USD 7.5 billion to USD 32.4 billion in a single quarter; even after this issuance, NVIDIA remained in a net cash position of approximately USD 66.0 billion. Working capital and receivables growth have kept pace with revenue growth without evident deterioration in disclosed inventory or receivables quality, though the company does not break out days-sales-outstanding or inventory-days metrics in its headline release, and a full aging analysis is Not Retrieved from the sources used for this report. No off-balance-sheet financing structures or pension obligations material to the analysis were identified in the sources reviewed.

C3. Cash Flow Statement

CASH FLOW SUMMARY, FY2022-FY2026 (US\$ MILLIONS)

Line item	FY2022	FY2023	FY2024	FY2025	FY2026
Operating cash flow	9,108	5,641	28,090	64,089	102,718
Capital expenditure	(976)	(1,833)	(1,069)	(3,236)	(6,042)
Free cash flow	8,132	3,808	27,021	60,853	96,676
FCF margin	30.2%	14.1%	44.4%	46.6%	44.8%
Share repurchases	n.a.	n.a.	n.a.	n.a.	(40,086)
Dividends paid	n.a.	n.a.	n.a.	n.a.	(974)

Sourced: NVIDIA fiscal-year results releases and stockanalysis.com (see Section G). FY2022-FY2025 buyback and dividend detail not separately retrieved and shown as not available rather than estimated. FCF margin is Reasoned/Calculated as free cash flow divided by revenue for the corresponding year.

Operating cash flow and free cash flow, FY2022-FY2026

Units: US\$ billions. Source: NVIDIA results releases and stockanalysis.com.

For the first half of fiscal 2027 (February to July 2026), NVIDIA reported operating cash flow of USD 74.4 billion and free cash flow of USD 69.9 billion, on capital expenditure of roughly USD 4.5 billion for the half; annualizing this run-rate would imply full-year free cash flow approaching USD 140 billion, though the company does not guide free cash flow directly and this report treats any annualization as Reasoned rather than as guidance. Capital allocation in FY2026 returned USD 41.1 billion to shareholders (USD 40.1 billion of buybacks plus USD 974 million of dividends) against free cash flow of USD 96.7 billion, a payout ratio of approximately 43 percent of free cash flow, leaving the majority of free cash flow available for reinvestment or balance-sheet strengthening. As flagged in Section B1, FY2026 operating cash flow (USD 102.7 billion) ran below net income (USD 120.1 billion); this is consistent with working-capital consumption (build-up of receivables and inventory to support triple-digit revenue growth and forward supply commitments) rather than an indication that reported earnings are not backed by cash, since free cash flow of USD 96.7 billion still represented conversion of more than 80 percent of net income into cash after capital spending.

C4. Forecasting and Projections

The starting point for any forward model is the trajectory already visible in actual results and guidance: first-half fiscal 2027 revenue of USD 177.8 billion (up 96 percent year over year) and third-quarter fiscal 2027 guidance of USD 108.0 billion, plus or minus 2 percent (up approximately 61 percent year over year at the midpoint), both Sourced from NVIDIA's own disclosures. The key open variable is the pace of deceleration from this base as year-over-year comparisons become progressively harder.

REVENUE SCENARIO FRAMEWORK, FY2027E-FY2031E (REASONED/ESTIMATED, US\$ BILLIONS)

Scenario	FY2027E	FY2028E	FY2029E	FY2030E	FY2031E	FY2032E
Bull (revenue growth 95%, 42%, 28%, 20%, 15%)	421	598	765	918	1,056	n.a.
Base (revenue growth 87%, 32%, 20%, 14%, 10%)	404	533	640	729	802	n.a.
Bear (revenue growth 70%, 15%, 8%, 5%, 4%)	367	422	456	479	498	n.a.

Reasoned/Estimated. FY2027E growth of 87 percent (Base) reflects actual first-half results plus the company's own Q3 guidance plus an assumed, unguided Q4 sequential increase; growth rates for FY2028E onward are analyst assumptions, not company guidance, and decelerate progressively as the revenue

base scales, consistent with the deceleration already visible in reported quarterly growth (from over 100 percent year-over-year in late FY2026 toward the 61 percent implied by Q3 FY2027 guidance).

Margin projections. Base case assumes free cash flow margin normalizing from the FY2026 level of 44.8 percent toward a still-high 38 percent by the fifth forecast year, reflecting an assumed gradual increase in capital intensity (own data-center and energy investment, debt service) and the historical tendency for hardware gross margins to compress modestly as competitive intensity in AI accelerators increases. Bull case assumes margins hold closer to current levels (40 to 44 percent) on continued pricing power; Bear case assumes faster margin compression (down to 30 percent) reflecting more aggressive competitive and customer-concentration pricing pressure. **Working capital and capital expenditure:** NVIDIA's own reported capital expenditure remains a small percentage of revenue (2.8 percent in FY2026) because manufacturing is outsourced; this report assumes capital intensity rises modestly but remains well below 10 percent of revenue across the forecast horizon, consistent with the fabless model, while noting the new USD 25.0 billion debt issuance suggests a larger share of future capital spending (including in NVIDIA's own infrastructure and energy commitments) may be debt-funded going forward.

Sensitivity. The valuation in Section C5 is most sensitive, in descending order, to (i) the assumed revenue growth path over the explicit forecast period, (ii) the terminal free cash flow margin, (iii) the weighted average cost of capital, and (iv) the terminal growth rate; the sensitivity table in Section C5 quantifies the effect of the latter two.

C5. Discounted Cash Flow Model

DCF ASSUMPTIONS (BASE CASE)

Assumption	Value	Basis
Base-year revenue (FY2026)	USD 215.9bn	Sourced
Explicit forecast horizon	5 years	Reasoned
Revenue growth path (Base)	87% / 32% / 20% / 14% / 10%	Reasoned, see Section C4
Free cash flow margin path (Base)	42% declining to 38%	Reasoned, see Section C4
Risk-free rate	4.68%	Sourced: 10-year U.S. Treasury yield, August 2026
Beta	2.22	Sourced/cross-validated, Section B4
Equity risk premium	5.00%	Reasoned: standard long-run U.S. equity risk premium assumption
Cost of equity (CAPM)	15.77%	Reasoned/Calculated: $4.68\% + 2.22 \times 5.00\%$
Pre-tax cost of debt	4.50%	Reasoned: estimate for a high-grade issuer given the September 2026 senior notes issuance; exact coupon Not Retrieved
Weighted average cost of capital	15.7%	Reasoned/Calculated, weighted at market value (approximately 99% equity / 1% debt)
Terminal growth rate (Base)	3.5%	Reasoned
Net cash (Jul-2026)	USD 66.0bn	Sourced

Assumption	Value	Basis
Diluted shares outstanding	24,147m	Sourced

DCF VALUATION RANGE BY SCENARIO

Scenario	WACC	Terminal growth	Implied enterprise value	Implied equity value per share
Bear	17.2%	2.5%	USD 0.93tn	USD 41
Base	15.7%	3.5%	USD 2.02tn	USD 86
Bull	14.7%	4.0%	USD 3.07tn	USD 130

Reasoned/Calculated by the analyst using the assumptions above; not a price target or a recommendation. The wide range (USD 41 to USD 130 per share) reflects the genuine sensitivity of a high-growth, high-multiple business to small changes in growth and margin assumptions over a long explicit forecast horizon, and is presented as a range rather than a point estimate as required by the report's methodology.

BASE-CASE SENSITIVITY: IMPLIED VALUE PER SHARE BY WACC AND TERMINAL GROWTH

WACC \ Terminal growth	3.0%	3.5%	4.0%
14.2%	USD 96	USD 99	USD 102
15.7%	USD 84	USD 86	USD 89
17.2%	USD 75	USD 77	USD 78

Reasoned/Calculated, holding the Base-case revenue and margin path constant and varying only the discount rate and terminal growth rate.

Every scenario in this report's DCF framework, including the Bull case, implies a per-share value materially below the 25 September 2026 closing price of USD 225.07. This does not mean the market is necessarily wrong: it means the market is pricing in a growth and margin path more aggressive than even this report's Bull case, over a longer horizon than the five years modeled explicitly, or it is applying a lower discount rate than the CAPM-derived 15.7 percent used here, for example by treating NVIDIA's cash flows as less risky than its historical beta would suggest. Readers should treat this gap as a quantified description of market expectations (see Section B3), not as a signal to act.

C6. Leveraged Buyout Model

A leveraged buyout analysis is not a meaningful valuation approach for NVIDIA and is provided in summary form only, consistent with the report's methodology of omitting analyses that are not economically meaningful for the company being analyzed. NVIDIA's market capitalization of approximately USD 5.4 trillion is orders of magnitude beyond the financing capacity of any realistic private-equity consortium or debt market; its net cash balance sheet and modest, recently introduced leverage (net cash of USD 66.0 billion even after a USD 25.0 billion bond issuance) is the opposite of the balance-sheet profile an LBO sponsor typically targets; and its value is overwhelmingly driven by continued heavy reinvestment in research and development and by a return profile (ROIC near 70 percent, Section B1) that would be difficult to leverage further without impairing the very growth that creates the value a sponsor would be underwriting. For these reasons this report does not construct a hypothetical debt paydown schedule or sponsor returns analysis for NVIDIA.

D. Investment Thesis

Bull Case

AI infrastructure capital spending among hyperscale and sovereign buyers continues to grow at a pace that fully absorbs NVIDIA's expanding accelerator, networking and system output, consistent with management's own commentary that supply, not demand, is presently the binding constraint. Data Center revenue continues to compound at a rate closer to this report's Bull-case growth path (95 percent in FY2027, moderating toward 15 percent by FY2032) than its Base case, aided by the Rubin architecture transition sustaining premium pricing and by system-level (rather than component-level) sales capturing a growing share of total AI infrastructure spending. Gross and free-cash-flow margins hold near current elevated levels rather than normalizing downward, as CUDA-based switching costs and full-stack system integration continue to support pricing power even as competitive intensity from custom silicon rises. Under these conditions, this report's own DCF framework (Section C5) would support a materially higher intrinsic value than the Base case, and continued share repurchases and an enlarged dividend would compound per-share value creation on top of underlying growth.

Bear Case

Growth decelerates faster than guided, either because one or more of the four customers that presently account for 61 percent of revenue slow capital spending, or because in-house silicon programmes at those same customers substitute for a meaningful share of merchant GPU purchases sooner than the market expects. Margins compress as competitive intensity increases from AMD, Broadcom-enabled custom silicon, and in-house accelerators, while the cost base normalizes as NVIDIA's own infrastructure and energy commitments (partly funded by the new USD 25.0 billion of debt) increase capital intensity. Further export-control tightening removes additional addressable demand beyond the China market already foreclosed as of fiscal year end 2026. Under this combination, this report's Bear-case DCF (WACC 17.2 percent, terminal growth 2.5 percent) implies a value near USD 41 per share, and the stock's beta of 2.2 and historical maximum drawdown of 62.8 percent illustrate how sharply a growth-and-multiple business of this kind can re-rate once growth expectations are revised downward.

Base Case

Revenue growth decelerates from the currently guided 61 percent (Q3 FY2027 implied year-over-year growth) toward the 30s percent range over the next one to two fiscal years and further toward low-double-digit growth by the early 2030s, broadly consistent with the deceleration path already visible in NVIDIA's own quarterly results. Free cash flow margin normalizes gradually from the current high-40s percent level toward the high-30s percent range as capital intensity rises and competitive pressure increases, while remaining well above typical semiconductor industry levels given the software and systems components of the business. Capital allocation continues to combine heavy R&D reinvestment, continued buybacks, a materially larger dividend than in prior years, and moderate, opportunistic use of debt. On these assumptions, this report's Base-case DCF (Section C5) implies a per-share value near USD 86, below the current market price, indicating the market is pricing in a more optimistic path than this Base case, as quantified in Section B3.

Catalysts

- Quarterly earnings releases, particularly Data Center revenue growth and gross margin trends relative to the Q3 FY2027 guidance of USD 108.0 billion revenue and 74.0 percent gross margin
- Rubin architecture product ramp and customer adoption, as the successor to the current Blackwell generation

- Further hyperscale capital-expenditure guidance from NVIDIA's largest customers, which are leading indicators for Data Center demand
- Changes to United States export-control policy toward China and other restricted markets
- Continued execution of the buyback programme and the newly enlarged dividend
- Progress of the disclosed Intel equity and collaboration relationship, and any similar strategic investments

Things to Monitor

MONITORING LIST

Metric or event	Why it matters	Positive signal	Negative signal
Data Center revenue growth (quarterly)	Determines the trajectory of consolidated results	Growth stable or above guidance	Growth decelerating faster than guided
Customer concentration disclosure	Quantifies dependence on a small buyer base	Concentration stable or declining	Concentration continuing to rise toward a majority of revenue from one or two customers
Gross margin	Signals pricing power versus competitive and cost pressure	Margin holding near or above 74 percent	Margin trending back toward the high-60s percent range outside of one-off charges
Hyperscaler capital-expenditure guidance	Leading indicator of NVIDIA Data Center demand	Capex guidance raised or maintained	Capex guidance cut or growth explicitly redirected toward in-house silicon
Net debt and bond issuance activity	Tracks whether September 2026 was a one-off or the start of a more leveraged capital structure	Net cash position maintained or rebuilt	Repeated large debt issuances shifting the company toward a net-debt position
Export-control and trade-policy developments	Can add or remove entire addressable markets	Market access stable or expanding	Further restrictions removing additional addressable revenue

E. Final Investment Framework

BUSINESS QUALITY	High. Dominant position in AI accelerator compute with a software ecosystem (CUDA) that provides durable switching costs; gross margin sustained above 70 percent through more than tenfold volume growth.
FINANCIAL QUALITY	High. Net cash balance sheet even after a USD 25.0 billion bond issuance, free cash flow conversion above 40 percent of revenue, ROIC near 70 percent.
COMPETITIVE POSITION	Strong but facing rising pressure. Leading merchant-silicon position; increasing exposure to hyperscaler in-house silicon and Broadcom-enabled custom accelerators among its own largest customers.
MANAGEMENT	Long-tenured founder-CEO with a multi-decade record of successful strategic pivots; capital allocation combines heavy reinvestment with substantial, increasing shareholder returns and only recently introduced leverage.
INDUSTRY	Growing and structurally transforming, with supply (not demand) the near-term binding constraint; narrow, capital-intensive customer base is the principal structural vulnerability.
VALUATION	Demanding on every method examined in Section B3; this report's own DCF framework implies a Base-case value materially below the current price, indicating the market is pricing in growth and margin assumptions above this report's Base case.
TECHNICAL AND MARKET STRUCTURE	Intact medium-term uptrend (price above both 50-day and 200-day moving averages) with a neutral RSI; elevated realized volatility (40 to 50 percent annualized) and beta (2.2) versus the broad market.
PRINCIPAL RISK	A structural shift by NVIDIA's largest customers toward in-house AI accelerators, compounded by single-foundry dependence on TSMC.
PRINCIPAL CATALYST	Continued execution of the Data Center revenue and margin trajectory implied by current guidance, alongside the Rubin architecture product ramp.
KEY VARIABLE	The rate of deceleration of Data Center revenue growth relative to the growth path already embedded in the current share price.
THESIS INVALIDATION	The investment thesis would become materially weaker if two or more of NVIDIA's largest customers publicly redirected a majority of incremental AI infrastructure capital expenditure toward in-house silicon, if gross margin fell and remained below 65 percent for more than two consecutive quarters outside of a disclosed one-off charge, or if the net cash position turned into sustained net debt without a corresponding acceleration in free cash flow.

F. Data Table

STANDARDIZED COMPANY SNAPSHOT

Category	Metric	Current	Prior Year	3Y / 5Y Trend
Market	Market Capitalization	USD 5,434.8bn	USD 4,634.4bn (FY2026-end)	Up from USD 610bn (FY2022-end)
Market	Enterprise Value	USD 5,368.8bn	USD 4,580.3bn (FY2026-end)	Not separately tracked pre-FY2026
Market	Share Price	USD 225.07	USD 190.68 (FY2026-end)	5-year: +991.2%
Growth	Revenue Growth (YoY)	40.3% (TTM)	65.5% (FY2026)	3-year CAGR: 100.2%
Profitability	Gross Margin	74.7% (TTM)	71.1% (FY2026)	FY2022: 64.9%
Profitability	Operating Margin	65.2% (TTM)	60.4% (FY2026)	FY2022: 37.3%
Profitability	Net Margin	63.7% (TTM)	55.6% (FY2026)	FY2022: 36.2%
Profitability	ROE	84.2% (TTM)	76.3% (FY2026)	FY2022: 36.6%
Profitability	ROIC	69.9% (TTM)	71.3% (FY2026)	Calculated, Section B1
Cash Flow	Operating Cash Flow	USD 74.4bn (H1 FY2027)	USD 102.7bn (FY2026)	FY2022: USD 9.1bn
Cash Flow	Free Cash Flow	USD 69.9bn (H1 FY2027)	USD 96.7bn (FY2026)	FY2022: USD 8.1bn
Leverage	Net Debt / EBITDA	Net cash (approx. -0.3x)	Net cash (FY2026)	Net cash throughout FY2022-Q2 FY2027
Valuation	P/E	28.5x	38.9x (FY2026-end)	FY2022-end: 62.5x
Valuation	EV/EBITDA	23.2x - 26.7x	Not directly comparable (EBITDA basis differs)	See Section B3
Valuation	P/FCF	38.9x (H1 FY2027 annualized)	56.2x (FY2026 basis)	See Section B3
Market	Beta	2.22	n.a.	5-year monthly, calculated
Market	Volatility (annualized)	40.5% (3-month)	49.7% (5-year)	Section B4
Market	Maximum Drawdown	(19.4%) (3-month)	(8.8%) (2026 YTD)	(62.8%) (5-year, Nov-2021 to Sep-2022)

Sourced and Reasoned/Calculated as detailed throughout Sections B and C; see Section G for underlying sources. "Current" for income-statement and cash-flow items uses the most recent complete reporting period available (trailing twelve months to 26 July 2026, or the first half of fiscal 2027 where a clean trailing figure could not be reconstructed from Sourced data without estimation).

G. Sources and Methodology

Primary Sources

NVIDIA Corporation. "NVIDIA Announces Financial Results for Fourth Quarter and Fiscal 2026." Results release, 25 February 2026. Used for FY2026 and Q4 FY2026 revenue, margin, income, cash flow, segment and capital-allocation data.

NVIDIA Corporation. "NVIDIA Announces Financial Results for Second Quarter Fiscal 2027." Results release, 26 August 2026. Used for Q2 FY2027 and first-half FY2027 revenue, margin, income, cash flow, balance-sheet and guidance data.

U.S. Securities and Exchange Commission (EDGAR). NVIDIA Corporation, Form 10-K for the fiscal year ended 25 January 2026 (filed 2026). Used for customer-concentration risk disclosure, China export-control and H2O inventory-charge disclosure, and named executive officer information.

U.S. Securities and Exchange Commission (EDGAR). NVIDIA Corporation, Form 10-Q for the quarterly period ended 26 July 2026. Used for the Q2 FY2027 condensed balance sheet (cash, marketable securities, debt, equity, shares outstanding).

U.S. Securities and Exchange Commission (EDGAR). NVIDIA Corporation, CFO Commentary on Second Quarter Fiscal 2027 Results, filed as an exhibit 26 August 2026. Used for the senior notes issuance amount, inventory, cash and investments detail, and operating-expense detail.

NVIDIA Corporation. Investor Relations, Financial Reports page (investor.nvidia.com). Used to confirm the primary source location and headline Q2 FY2027 figures.

NVIDIA Corporation. Data Center product platform page (nvidia.com). Used for product and architecture description (Blackwell, Rubin, Grace, Vera, BlueField, Spectrum-X) in Section A2.

Secondary Sources

The Motley Fool. "Blackwell Sales Are Off the Charts for Nvidia, and Worryingly, So Is Its Customer Concentration." 27 November 2025. Used for the quarterly customer-concentration percentages cited in Sections A3, A5, B2, D and E.

Market Data

Alpha Vantage. Company Overview and daily/monthly time-series data for NVDA, AMD, AVGO and INTC; U.S. 10-year Treasury yield series. Accessed 26-27 September 2026. Used for current price, 52-week range, moving averages, beta, valuation multiples, peer financial data, and the DCF risk-free rate.

stockanalysis.com. NVIDIA Corporation (NVDA) financial statements (income statement, balance sheet, cash flow statement) and Taiwan Semiconductor Manufacturing Company (TSM) key statistics. Accessed 26-27 September 2026. Used for historical (FY2022-FY2026) financial-statement detail and for TSMC peer valuation data.

Calculations

Analyst calculations (Arthur Faugeron). All margin, growth-rate, return-on-capital, enterprise-value, valuation-multiple, beta, volatility, maximum-drawdown, RSI, discounted-cash-flow and sensitivity figures not directly disclosed by the company are Reasoned or Calculated by the analyst from the Sourced inputs listed above, using the formulas and assumptions stated at each point in the report where they appear, principally in Sections B1, B3, B4, C4 and C5.

This report has been prepared for personal research purposes only and does not constitute investment advice, financial advice, or a recommendation to buy, sell, or hold any security. The information contained herein has been compiled with the assistance of artificial intelligence and is based on publicly available sources believed to be reliable, but its accuracy and completeness are not guaranteed. Past performance is not indicative of future results. Investors should conduct their own due diligence and consult with a qualified financial professional before making any investment decisions.